

Project: Probability of Default (PD) Prediction Model with MLflow Tracking and Registry

This document provides a clear, beginner-friendly explanation of a machine learning (ML) mini-project using MLflow. It walks through the entire process step by step, even for those who have not used MLflow or built ML models before.

Objective

To build a machine learning model that predicts whether a customer will default on a loan (Probability of Default or PD), and to track the entire model development and deployment process using **MLflow**.

What is MLflow?

MLflow is a tool that helps you manage the lifecycle of machine learning models. With MLflow, you can:

- Track model parameters, metrics, and artifacts (like plots)
 - Save and load models.
 - Register models in a model registry.
 - Manage different versions of models.
 - Transition models to various stages (e.g., Staging, Production)
-

Step-by-Step Explanation of the Project

Step 1: Install Libraries

```
!pip install pandas scikit-learn matplotlib seaborn mlflow
```

These libraries help us with data processing (pandas), machine learning (scikit-learn), visualizations (matplotlib, seaborn), and MLflow tracking.

✅ Step 2: Import Libraries

We load all the required libraries for model training and tracking.

✅ Step 3: Generate Sample Data

```
from sklearn.datasets import make_classification
```

We generate a fake dataset that simulates customer data for loan applications, with a target variable representing whether they defaulted or not.

✅ Step 4: Train a Classification Model

```
from sklearn.ensemble import RandomForestClassifier
```

We train a **Random Forest classifier** to predict loan default based on customer features. We evaluate it using two metrics:

- **Accuracy:** How many predictions were correct
 - **ROC AUC:** How well the model separates defaulters from non-defaulters
-

✅ Step 5: Track the Model Using MLflow

```
mlflow.set_experiment("PD_Classification_Experiment")
```

```
with mlflow.start_run():
```

```
...
```

Inside this block, we:

- Log **parameters** (e.g., number of trees, depth)
 - Log **metrics** (e.g., accuracy, ROC AUC)
 - Log **artifacts** (e.g., confusion matrix plot, ROC curve)
 - Log and **register** the trained model to the **Model Registry**
-

✅ Step 6: Visualize Model Performance

We generate:

- A **confusion matrix**: Shows true/false predictions
 - A **ROC curve**: Plots model's performance across thresholds Both are logged as images in MLflow.
-

✅ Step 7: Register and Version the Model

```
mlflow.sklearn.log_model(model, "model",  
registered_model_name="pd_model_v1")
```

This saves the model in MLflow and adds it to a **named model registry** (pd_model_v1). Each time you log a new model with this name, MLflow creates a new **version**.

✅ Step 8: Load the Model Later for Use

```
mlflow.sklearn.load_model("models:/pd_model_v1/latest")
```

You can load the saved model any time in the future and use it to make predictions.

✅ Step 9: Promote the Model to Production

```
mlflowClient().transition_model_version_stage(...)
```

This marks a version of the model as "**Production**", meaning it's now the official model to use. You can also transition models to:

- **Staging** (for testing)
 - **Archived** (for old versions)
-

✅ Step 10: Fine-Tune the Model

```
from sklearn.model_selection import GridSearchCV
```

We use grid search to find the best hyperparameters for our model and improve performance. The best model can also be logged and promoted using MLflow.

MLflow Dashboard

After starting the MLflow UI with:

```
mlflow ui
```

You can view all your:

- Experiments
- Runs
- Metrics & parameters
- Model versions
- Visual artifacts
- Registered models

Access it via: <http://127.0.0.1:5000>

Final Result

You've created a complete end-to-end ML pipeline that:

- Trains and evaluates a model.
- Logs everything automatically
- Registers and version-controls the model.
- Loads and reuses the model.
- Promotes the best model to production.

Why This Is Valuable

This mimics a real-world ML workflow used in companies:

- Track experiments
 - Compare results.
 - Reuse & deploy reliable models.
 - Keep everything reproducible & organized.
-